

# Assignment 2

April 12, 2020

---

You are currently looking at **version 1.2** of this notebook. To download notebooks and datafiles, as well as get help on Jupyter notebooks in the Coursera platform, visit the [Jupyter Notebook FAQ](#) course resource.

---

## 1 Assignment 2 - Pandas Introduction

All questions are weighted the same in this assignment. ## Part 1 The following code loads the olympics dataset (olympics.csv), which was derived from the Wikipedia entry on [All Time Olympic Games Medals](#), and does some basic data cleaning.

The columns are organized as # of Summer games, Summer medals, # of Winter games, Winter medals, total # number of games, total # of medals. Use this dataset to answer the questions below.

```
In [65]: import pandas as pd
```

```
df = pd.read_csv('olympics.csv', index_col=0, skiprows=1)

for col in df.columns:
    if col[:2]=='01':
        df.rename(columns={col:'Gold'+col[4:]}, inplace=True)
    if col[:2]=='02':
        df.rename(columns={col:'Silver'+col[4:]}, inplace=True)
    if col[:2]=='03':
        df.rename(columns={col:'Bronze'+col[4:]}, inplace=True)
    if col[:1]=='':
        df.rename(columns={col:'#'+col[1:]}, inplace=True)

names_ids = df.index.str.split('\s\(')
df.index = names_ids.str[0]

df = df.drop('Totals')
df.head()
```

```
Out[65]:
```

	# Summer	Gold	Silver	Bronze	Total	# Winter	Gold.1	\
Afghanistan	13	0	0	2	2	0	0	

Algeria	12	5	2	8	15	3	0
Argentina	23	18	24	28	70	18	0
Armenia	5	1	2	9	12	6	0
Australasia	2	3	4	5	12	0	0

	Silver.1	Bronze.1	Total.1	# Games	Gold.2	Silver.2	Bronze.2	\
Afghanistan	0	0	0	13	0	0	2	
Algeria	0	0	0	15	5	2	8	
Argentina	0	0	0	41	18	24	28	
Armenia	0	0	0	11	1	2	9	
Australasia	0	0	0	2	3	4	5	

	Combined total
Afghanistan	2
Algeria	15
Argentina	70
Armenia	12
Australasia	12

### 1.0.1 Question 0 (Example)

What is the first country in df?

*This function should return a Series.*

```
In [66]: # You should write your whole answer within the function provided. The autograder will
# this function and compare the return value against the correct solution value
def answer_zero():
    # This function returns the row for Afghanistan, which is a Series object. The assignment
    # question description will tell you the general format the autograder is expecting
    return df.iloc[0]

# You can examine what your function returns by calling it in the cell. If you have questions
# about the assignment formats, check out the discussion forums for any FAQs
answer_zero().name
```

Out[66]: 'Afghanistan'

### 1.0.2 Question 1

Which country has won the most gold medals in summer games?

*This function should return a single string value.*

```
In [67]: def answer_one():
    return df['Gold'].argmax()
answer_one()
```

Out[67]: 'United States'

### 1.0.3 Question 2

Which country had the biggest difference between their summer and winter gold medal counts?  
*This function should return a single string value.*

```
In [68]: def answer_two():  
         return (df['Gold'] - df['Gold.1']).idxmax()  
         answer_two()
```

```
Out[68]: 'United States'
```

### 1.0.4 Question 3

Which country has the biggest difference between their summer gold medal counts and winter gold medal counts relative to their total gold medal count?

$$\frac{\text{Summer Gold} - \text{Winter Gold}}{\text{Total Gold}}$$

Only include countries that have won at least 1 gold in both summer and winter.  
*This function should return a single string value.*

```
In [69]: def answer_three():  
         only_gold = df.where((df['Gold'] > 0) & (df['Gold.1'] > 0))  
         only_gold = only_gold.dropna()  
         return (abs((only_gold['Gold'] - only_gold['Gold.1']) / only_gold['Gold.2'])).idxmax()  
  
         answer_three()
```

```
Out[69]: 'Bulgaria'
```

### 1.0.5 Question 4

Write a function that creates a Series called "Points" which is a weighted value where each gold medal (Gold.2) counts for 3 points, silver medals (Silver.2) for 2 points, and bronze medals (Bronze.2) for 1 point. The function should return only the column (a Series object) which you created, with the country names as indices.

*This function should return a Series named Points of length 146*

```
In [70]: def answer_four():  
         df['Points'] = (df['Gold.2'] * 3 + df['Silver.2'] * 2 + df['Bronze.2'] * 1)  
         return df['Points']  
  
         answer_four()
```

```
Out[70]: Afghanistan      2  
         Algeria          27  
         Argentina       130  
         Armenia          16  
         Australasia      22  
         Australia       923
```

Austria	569
Azerbaijan	43
Bahamas	24
Bahrain	1
Barbados	1
Belarus	154
Belgium	276
Bermuda	1
Bohemia	5
Botswana	2
Brazil	184
British West Indies	2
Bulgaria	411
Burundi	3
Cameroon	12
Canada	846
Chile	24
China	1120
Colombia	29
Costa Rica	7
Ivory Coast	2
Croatia	67
Cuba	420
Cyprus	2
...	...
Spain	268
Sri Lanka	4
Sudan	2
Suriname	4
Sweden	1217
Switzerland	630
Syria	6
Chinese Taipei	32
Tajikistan	4
Tanzania	4
Thailand	44
Togo	1
Tonga	2
Trinidad and Tobago	27
Tunisia	19
Turkey	191
Uganda	14
Ukraine	220
United Arab Emirates	3
United States	5684
Uruguay	16
Uzbekistan	38
Venezuela	18

Vietnam	4
Virgin Islands	2
Yugoslavia	171
Independent Olympic Participants	4
Zambia	3
Zimbabwe	18
Mixed team	38

Name: Points, dtype: int64

## 1.1 Part 2

For the next set of questions, we will be using census data from the [United States Census Bureau](#). Counties are political and geographic subdivisions of states in the United States. This dataset contains population data for counties and states in the US from 2010 to 2015. [See this document](#) for a description of the variable names.

The census dataset (census.csv) should be loaded as census\_df. Answer questions using this as appropriate.

### 1.1.1 Question 5

Which state has the most counties in it? (hint: consider the sumlevel key carefully! You'll need this for future questions too...)

*This function should return a single string value.*

```
In [35]: census_df = pd.read_csv('census.csv')
        census_df.head()
```

```
Out[35]:
```

	SUMLEV	REGION	DIVISION	STATE	COUNTY	STNAME	CTYNAME	\
0	40	3	6	1	0	Alabama	Alabama	
1	50	3	6	1	1	Alabama	Autauga County	
2	50	3	6	1	3	Alabama	Baldwin County	
3	50	3	6	1	5	Alabama	Barbour County	
4	50	3	6	1	7	Alabama	Bibb County	

  

	CENSUS2010POP	ESTIMATESBASE2010	POPESTIMATE2010	...	\
0	4779736	4780127	4785161	...	
1	54571	54571	54660	...	
2	182265	182265	183193	...	
3	27457	27457	27341	...	
4	22915	22919	22861	...	

  

	RDOMESTICMIG2011	RDOMESTICMIG2012	RDOMESTICMIG2013	RDOMESTICMIG2014	\
0	0.002295	-0.193196	0.381066	0.582002	
1	7.242091	-2.915927	-3.012349	2.265971	
2	14.832960	17.647293	21.845705	19.243287	
3	-4.728132	-2.500690	-7.056824	-3.904217	
4	-5.527043	-5.068871	-6.201001	-0.177537	

	RDOMESTICMIG2015	RNETMIG2011	RNETMIG2012	RNETMIG2013	RNETMIG2014	\
0	-0.467369	1.030015	0.826644	1.383282	1.724718	
1	-2.530799	7.606016	-2.626146	-2.722002	2.592270	
2	17.197872	15.844176	18.559627	22.727626	20.317142	
3	-10.543299	-4.874741	-2.758113	-7.167664	-3.978583	
4	0.177258	-5.088389	-4.363636	-5.403729	0.754533	

	RNETMIG2015
0	0.712594
1	-2.187333
2	18.293499
3	-10.543299
4	1.107861

[5 rows x 100 columns]

```
In [71]: def answer_five():
         new_df = census_df[census_df['SUMLEV'] == 50]
         return new_df.groupby('STNAME').count()['SUMLEV'].idxmax()

         answer_five()
```

Out[71]: 'Texas'

### 1.1.2 Question 6

**Only looking at the three most populous counties for each state**, what are the three most populous states (in order of highest population to lowest population)? Use CENSUS2010POP.

*This function should return a list of string values.*

```
In [72]: def answer_six():
         new_df = census_df[census_df['SUMLEV'] == 50]
         most_populous_counties = new_df.sort_values('CENSUS2010POP', ascending=False).groupby('STNAME').sum().sort_values('CENSUS2010POP', ascending=False)
         return most_populous_counties.groupby('STNAME').sum().sort_values('CENSUS2010POP', ascending=False)

         answer_six()
```

Out[72]: ['California', 'Texas', 'Illinois']

### 1.1.3 Question 7

Which county has had the largest absolute change in population within the period 2010-2015? (Hint: population values are stored in columns POPESTIMATE2010 through POPESTIMATE2015, you need to consider all six columns.)

e.g. If County Population in the 5 year period is 100, 120, 80, 105, 100, 130, then its largest change in the period would be  $|130-80| = 50$ .

*This function should return a single string value.*

```
In [73]: def answer_seven():
    new_df1 = census_df[census_df['SUMLEV'] == 50][[6, 9, 10, 11, 12, 13, 14]]
    new_df1["MaxDiff"] = abs(new_df1.max(axis=1) - new_df1.min(axis=1))
    most_change = new_df1.sort_values(by=["MaxDiff"], ascending = False)
    return most_change.iloc[0][0]

    answer_seven()

Out[73]: 'Harris County'
```

#### 1.1.4 Question 8

In this datafile, the United States is broken up into four regions using the "REGION" column.

Create a query that finds the counties that belong to regions 1 or 2, whose name starts with 'Washington', and whose POPESTIMATE2015 was greater than their POPESTIMATE 2014.

*This function should return a 5x2 DataFrame with the columns = ['STNAME', 'CTYNAME'] and the same index ID as the census\_df (sorted ascending by index).*

```
In [11]: def answer_eight():
    return "YOUR ANSWER HERE"
```